AGENTIC
LEARNING LABS

AI Prompt Playbook

12 prompts you can use this week.
Copy. Paste. Edit. Done.

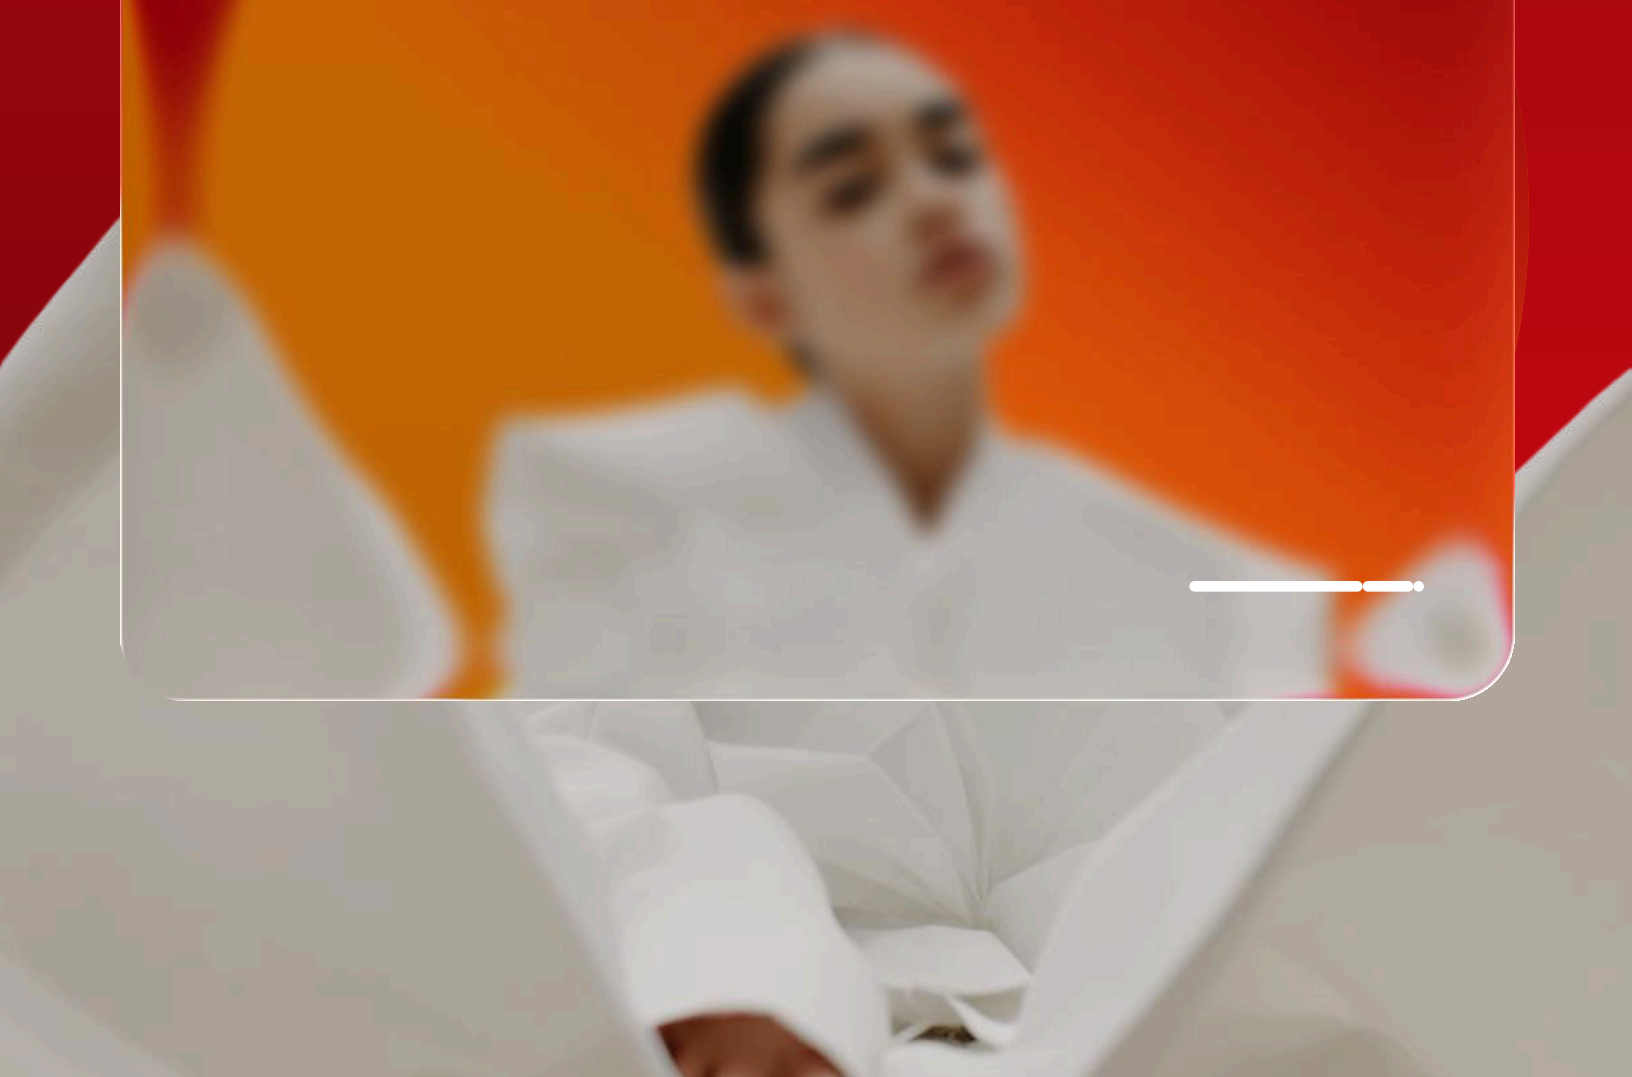

BEFORE YOU START

Most people use AI like a search engine. They type a question, get an answer, and wonder why it keeps missing the mark.

AI works best when you treat it like a smart intern who needs clear instructions. Give it context. Tell it what you want. Check the result.

The prompts in this playbook help give the AI the context it needs to produce something useful on the first try, and build in checkpoints so you catch problems before they become rewrites

There are 12 prompts organized into four sections:

- Act I covers the basics that make everything else work better.
- Act II helps you see your work through your audience's eyes before they do.
- Act III moves from describing what you want to actually building it.
- Act IV sets up agents that do recurring work for you.

You do not need a technical background to use any of these. Each one includes the prompt itself, what you will get back, a pro move to get more out of it, and a watch-out for the most common mistakes.

ACT I - GET STARTED

Prompts 1-3: The basics that make everything else work better

1 - Brief Your AI Like a New Hire

Time saved: 15-30 min of back-and-forth → under 5 min to a usable first draft

The situation: You need a first draft - an email, a summary, a plan - and you're not sure how to start.

The Prompt: You are helping me with a work task. Here is everything you need to know:
Context: [What's going on? 2-3 sentences of background.]
Audience: [Who is this for? Their role, what they care about.]
Task: [What do I need you to do? Be specific.]
Format: [How should the output look? Bullets, paragraphs, a table?]
Length: [How long should it be?]
Do not: [Anything to avoid - jargon, a certain tone, invented facts.]

Before you start, tell me in two sentences what you think I'm asking for.

What you'll get:

- A confirmation of the task before AI starts
- A first draft that matches your format and audience
- Flagged gaps if your input was missing something

✅ **Pro Move:** If the two-sentence confirmation is wrong - correct that, not the whole prompt. One fix saves five rounds of rewriting.

⚠️ **Watch Out:** If you skip the "Do not" line, AI will make assumptions. It always fills gaps - usually with something generic.

2 - The Confirmation Check

Time saved: Prevents 1-3 rounds of "that's not what I meant" rewrites

The situation: You paste a prompt and get back something that looks right but misses the point. This one-line addition stops that from happening. Add it to any prompt, any time.

The Prompt: Before you write anything, read my request and tell me:
1. What you think I'm asking for
2. Who the audience is
3. Any information you're missing

Wait for me to confirm before you start.

Add this to the top or bottom of any prompt you already use.

What you'll get:

- A short summary of what AI understood
- A list of gaps it needs you to fill
- A better output on the first try

✅ **Pro Move:** When AI lists gaps, answer only those. Don't rewrite your whole prompt. Targeted answers are faster.

⚠️ **Watch Out:** Skipping this step is the number one reason people get polished outputs that are completely wrong.

3 - The Self-Check

Time saved: Replaces a full review cycle with a 60-second AI audit

The situation: AI is confident whether it's right or wrong. This prompt makes it check its own work before you do - and flag problems for you.

The Prompt: Review the output you just gave me. Score it 1-5 on each of these:
- Accuracy: Is everything here actually true based on what I told you?
- Completeness: Did you cover everything I asked for?
- Tone: Does it sound right for the audience?
- Risk: Is there anything here I shouldn't send without checking?

For any score below 4, tell me exactly what the problem is and fix it.

What you'll get:

- A scored review of the draft
- Specific problems called out - not just "it could be better"
- A revised version with fixes applied

✅ **Pro Move:** Add "List three things a skeptical reader would question." This catches gaps you didn't know were there.

⚠️ **Watch Out:** This speeds up your review - it doesn't replace it. Always read the final version yourself before it goes anywhere.

ACT II - UNDERSTAND YOUR AUDIENCE

Prompts 4-6: See your work through someone else's eyes before they do

4 - The Customer Lens Test

Time saved: Replaces a focus group or days of guessing - feedback in 2 minutes

The situation: You've written a message, offer, or web page. You think it's clear. But will your customer? This prompt simulates how a specific type of person would read it - and where they'd walk away.

The Prompt: Read the content below as if you are: [describe the person - their role, what they care about, what frustrates them].

After reading, tell me:
1. What's your first reaction?
2. What's unclear or confusing?
3. What's missing that you'd want to know?
4. Would you take action - yes or no, and why?

Be direct. Don't be polite. A real customer wouldn't be.

[Paste your content here]

What you'll get:

- A persona-voiced reaction - not generic feedback
- Specific lines or sections that would lose the reader
- A clear yes/no on whether it would work

✅ **Pro Move:** Run the same content through two different customer types and compare. A price-sensitive buyer and a quality-first buyer will often react completely differently.

⚠️ **Watch Out:** This is a filter, not a final answer.

5 - The Rubric Scorecard

Time saved: Turns "this feels off" into a specific, fixable score in under 3 minutes

The situation: You're reviewing a proposal, job post, pitch, or piece of marketing. "Feels right" isn't a standard. This prompt makes AI give you a real score - and demand improvements, not compliments.

The Prompt: Score the content below on a scale of 1-5 for each criterion:
- Clarity: Is it easy to understand in one read?
- Audience fit: Does it match what [audience] actually cares about?
- Credibility: Does it back up claims, or just assert them?
- Action: Does it make clear what to do next?

For any score below 4, give me one specific fix - not general advice.
After fixes, show me the improved version.

[Paste your content here]

What you'll get:

- A scored table with one fix per weak area
- A revised version that hits your standards
- A repeatable benchmark you can use every time

✅ **Pro Move:** Set a rule: nothing goes out with any score below 4. AI will enforce your standard if you tell it to.

⚠️ **Watch Out:** If you don't define what a "5" looks like, AI will be generous. Add one sentence describing your ideal for each criterion.

6 - The Skeptic Test

Time saved: Surfaces the hardest objections before your meeting - not during it

The situation: You're about to present something important - a proposal, a decision, a change. Someone in the room will push back hard. This prompt finds those objections first so you're ready.

The Prompt: You are a [skeptical CFO / cautious investor / critical board member - pick one] who has seen a lot of overconfident proposals.

Read what I've written below and give me:
1. The top 5 objections you'd raise
2. The one "fatal flaw" that would kill the deal for you
3. A suggested response for each objection

Be tough. Generic pushback isn't useful.

[Paste your proposal, plan, or announcement here]

What you'll get:

- The five hardest questions you'll face - before the meeting
- A fatal flaw flag if one exists
- Suggested responses you can actually use

✅ **Pro Move:** After the skeptic review, ask AI to argue *in favor* of your proposal. See how strong the case actually is. If it's weak, rethink before you present.

⚠️ **Watch Out:** A generic skeptic gives generic pushback. The more specific your persona - industry, role, concerns - the more useful the output.

ACT III - BUILD AND VISUALIZE

Prompts 7-9: Stop describing what you want. Build it and see it instead.

7 - The 5-Minute Summary

Time saved: 20-minute read → 5-minute brief with higher recall

The situation: Someone sends you a long report, proposal, or article. You don't have time to read all of it. You need the key points - fast - in a format you'll actually remember.

The Prompt: I'm going to paste a document below. Summarize it for me as follows:

- The point in one sentence
- 5 key takeaways (one sentence each, plain language)
- What I need to decide or act on
- What's missing or unclear
- One question I should ask after reading this

Do not use jargon. Write like you're explaining to someone who hasn't read it.

[Paste document here]

What you'll get:

- A one-sentence headline
- 5 clear takeaways you can recall without rereading
- A decision or action prompt
- Gaps and questions called out

✅ **Pro Move:** Add "Flag any numbers, dates, or claims I should verify before sharing this." Saves you from forwarding something with an error in it.

⚠️ **Watch Out:** AI summarizes what's there - not what's missing. Always ask it to flag gaps explicitly.

8 - The Visual Dashboard

Time saved: Hours building slides or spreadsheets → a working visual in under 10 minutes

The situation: You have numbers - revenue, tasks, leads, metrics - and you want a clean visual to share with your team or review in a meeting. This prompt builds a working interactive dashboard. No design skills needed.

The Prompt: Build me an interactive HTML dashboard using the data below.

Include:
- A card for each metric showing the current value
- A color indicator: green = on track, yellow = watch, red = off track
- A trend arrow if I've given you more than one period of data
- A "Needs Attention" section for anything in red

Make it clean and easy to read on a laptop screen.

[Paste your metrics here - label each one and include targets if you have them]

What you'll get:

- A working HTML file you can open in any browser
- Color-coded metric cards with status indicators
- A "Needs Attention" section that pulls problems automatically

✅ **Pro Move:** Add "Include a one-line commentary under each metric explaining the status." Turns a visual into a briefing.

⚠️ **Watch Out:** The dashboard shows what you give it. AI doesn't connect to live data. Update the numbers manually when things change.

9 - The Rapid Prototype

Time saved: Weeks of planning → a working mockup in 15 minutes

The situation: You have an idea - a form, a calculator, a tool for your team or customers. Before you build it properly, you want to see if it actually works. This prompt creates a real, clickable version you can test today.

The Prompt: Build a working HTML prototype of the following tool:
What it does: [Describe the tool in 2-3 sentences]
Who uses it: [Who is the end user?]
Key inputs: [What does the user enter or select?]
Key output: [What should they see as a result?]

Keep it simple - this is for testing, not for launch.

After building it, list the three biggest assumptions you made
that I should check before going further.

What you'll get:

- A functional HTML prototype that works in a browser
- Placeholder data and working logic
- A list of assumptions to validate before investing more time

✅ **Pro Move:** Share the prototype with one real user before refining it. What they click first tells you more than any planning doc.

⚠️ **Watch Out:** Prototypes are for learning, not shipping. Treat them as disposable. The goal is to find out what's wrong before it's expensive to fix.

ACT IV - BUILD AGENTS THAT WORK FOR YOU

Prompts 10-12: Set up AI once. Use it every time.

10 - The Go / No-Go Decision Agent

Time saved: Hours of scattered research and gut-feel → a structured recommendation in one session

The situation: You're evaluating something big - a new product, a new hire, a new market, a new tool. You want a consistent, thorough process every time - not a different approach depending on the day. This agent walks you through it step by step.

The System Prompt *(paste this to start a new chat session):*

You are a structured decision advisor. Your job is to help me evaluate
whether to move forward with an idea.

Work through these phases in order. Do not skip ahead. After each phase,
summarize what you found and ask me to confirm before moving on.

Phase 1 - Understand the idea
Ask me 3-5 questions to understand what I'm evaluating and why.

Phase 2 - Analyze the opportunity
Based on my answers, outline: the potential upside, who it's for, and why now.

Phase 3 - Identify the risks
List the top 3 risks. For each one, rate it Low / Medium / High and explain why.

Phase 4 - Challenge the idea
Argue against moving forward. What would a cautious, experienced person say?

Phase 5 - Deliver a recommendation
Give me a structured output:

- One-line verdict: Go / No-Go / More Information Needed
- Top 3 reasons for your verdict
- The one thing I must resolve before deciding
- The decision and next step - that's mine to own

Rules:
- Do not invent facts. If you don't know something, say so.
- Flag anything I should verify with a real expert.
- Keep each phase under 200 words.

Start with Phase 1.

What you'll get:

- A guided conversation - not a wall of text
- A phase-by-phase analysis you can pause and add to
- A structured Go / No-Go report with a clear recommendation

✅ **Pro Move:** At the end, add: "Now argue the opposite of your recommendation." If the counter-argument is strong, you need more information before deciding.

⚠️ **Watch Out:** Thin input = thin output. The more specific you are in Phase 1, the more useful the recommendation. Don't rush through the questions.

11 - The Weekly Research Agent

Time saved: 1-2 hours of inconsistent searching → a structured research brief in 20 minutes

The situation: You or your team do the same type of research regularly - competitor tracking, supplier checks, market scans. The problem: everyone does it differently and you can't compare results. This agent standardizes it.

The System Prompt *(paste this to start a new chat session):*

You are a structured research assistant. Your job is to research a topic
and produce a consistent, useful report every time.

When I give you a topic, follow this process:

Step 1 - Clarify
Ask me two questions: What decision is this research supporting?
What would make this research useless?

Step 2 - Research the topic across these angles:
- What's the current situation?
- Who are the key players or options?
- What are the risks or downsides?
- What's the opportunity?
- What's changing or trending?

Step 3 - Rate your confidence
For each angle, rate your confidence 1-5 and explain why.
Flag anything you couldn't find.

Step 4 - Output the report
Use these sections:
Summary / Key Findings / Risks / Opportunity /
What I Couldn't Find / Suggested Next Steps

Rules:

- Do not invent facts or statistics.
- If something is uncertain, say so clearly.
- Keep the full report under 500 words.

Start by asking me what topic to research.

What you'll get:

- A consistent research format every time - no matter who runs it
- A confidence rating per section so you know what to verify
- A "What I Couldn't Find" flag - so gaps are visible, not hidden

✅ **Pro Move:** Save this prompt in a shared doc and give it to your team. Everyone gets the same quality of research output.

⚠️ **Watch Out:** AI finds what's easiest to find - not always what's most important. Define the angles that matter for your decision or it defaults to the obvious.

12 - Your Custom AI Specialist

Time saved: Replaces inconsistent judgment calls with a repeatable, reliable process

The situation: You have a task that needs consistent judgment - screening applicants, qualifying leads, reviewing proposals, responding to customers. Right now the quality depends on who's doing it. This agent makes it consistent every time, for everyone on your team.

The System Prompt Template *(fill in the blanks, then paste as your system prompt):*

You are [NAME], a specialist in [TASK].

Your job: [One sentence - what do you produce and who do you serve?]

How you work:
1. Start by asking for the input you need: [list 2-3 things]
2. Analyze it using these criteria: [list your criteria]
3. Produce a structured output (see format below)
4. Score yourself 1-5 on: accuracy, completeness, tone

5. Flag anything a human should review before acting on this

Output format (use every time):
- Summary: [what you found in 2-3 sentences]
- Assessment: [your structured evaluation]
- Flags: [anything uncertain, missing, or high-risk]
- Recommendation: [clear and specific]
- Human review required: Yes / No - and why
Rules:
- Do not invent facts or fill gaps with assumptions
- If information is missing, ask - don't guess
- Escalate to a human when: [list your 2-3 triggers]

Start by asking for the input you need.

What you'll get:

- A reusable agent you paste once and use going forward
- Consistent output format - every response looks the same
- Built-in escalation rules so you know when a human takes over

✅ **Pro Move:** Run three test scenarios before you deploy: a normal case, an edge case, and a sensitive one. Fix anything that scores below 4.

⚠️ **Watch Out:** Version 1 will not be perfect. Deploy it, collect the cases where it got something wrong, and update the prompt. The best agents improve with use.

WHAT'S NEXT?

These prompts are a starting point. The people who get the most out of AI, are the ones who take the time to learn the process, try new models and sharing learnings with others.

When one person on a team uses AI well, you get individual gains. When a whole team works this way, the compounding effect is something different entirely.

This is where Agentic Learning Labs comes in to help.

We deliver applied, human-centered AI training that workers can integrate into daily workflows from day one — no coding required. We teach AI as a tool of the job, not a replacement, through storytelling, hands-on practice, and proven change management.

Be the first to know: If you want to take your AI skills to the next level, join us for one of our new AI training programs coming April 2026. Visit www.agenticlearninglabs.com and sign up to be the first to hear when they launch!

Book a discovery call: hello@agenticlearninglabs.com

Let's chat to learn more about your organization's situation and how Agentic Learning Labs can help.

Or, follow us on social media:

@agenticlabs

Agentic Learning Labs

Quick Reference

#	Card	Best For	Time Saved
1	Brief Your AI Like a New Hire	Any first draft	15-30 min
2	The Confirmation Check	Preventing wrong outputs	1-3 rewrites
3	The Self-Check	Reviewing AI's own work	Full review cycle
4	The Customer Lens Test	Testing messages and offers	Focus group time
5	The Rubric Scorecard	Reviewing any content	Guesswork
6	The Skeptic Test	Preparing for tough meetings	Surprise objections
7	The 5-Minute Summary	Reading long documents	15+ min per doc
8	The Visual Dashboard	Sharing numbers visually	Hours of slide building
9	The Rapid Prototype	Testing ideas before building	Weeks of planning
10	The Go / No-Go Agent	Big decisions	Hours of scattered research
11	The Weekly Research Agent	Recurring research tasks	1-2 hr per cycle
12	Your Custom AI Specialist	Any repeatable judgment task	Ongoing inconsistency